

Heart Attack Prediction - Data Cleaning & EDA

1. Imports and Setup

```
In [7]: # Imports

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os

# Setup

os.makedirs('../data/processed', exist_ok=True)
os.makedirs('../reports/figures', exist_ok=True)
sns.set_style("whitegrid")
```

2. Load and Inspect Data

```
In [4]: df = pd.read_excel('../data/raw/data.xlsx', engine='openpyxl')

print("--- Data Info ---")
df.info()

print("\n--- Missing Values ---")
print(df.isnull().sum())

print(f"\n--- Number of Duplicate Rows ---")
print(df.duplicated().sum())
```

```
--- Data Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
--- Missing Values ---
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
--- Number of Duplicate Rows ---
```

```
1
```

3. Clean Data

```
In [5]: # --- Clean Data ---  
# Remove the 1 duplicate row found during inspection  
if df.duplicated().sum() > 0:  
    df.drop_duplicates(inplace=True)  
    print(f"Dropped {df.duplicated().sum() + 1} duplicate row. New shape: {df.shape}")  
else:  
    print("No duplicate rows found.")  
  
# Save the cleaned data for future use and for Tableau  
df.to_csv('../data/processed/cleaned_heart_data.csv', index=False)  
print("Cleaned data saved to 'data/processed/cleaned_heart_data.csv'")
```

Dropped 1 duplicate row. New shape: (302, 14)

Cleaned data saved to 'data/processed/cleaned_heart_data.csv'

4: Statistical Summary

```
In [8]: # --- Statistical Summary ---  
print("--- Statistical Summary of the Data ---")  
df.describe().T
```

```
--- Statistical Summary of the Data ---
```

Out[8]:

	count	mean	std	min	25%	50%	75%	max
age	302.0	54.420530	9.047970	29.0	48.00	55.5	61.00	77.0
sex	302.0	0.682119	0.466426	0.0	0.00	1.0	1.00	1.0
cp	302.0	0.963576	1.032044	0.0	0.00	1.0	2.00	3.0
trestbps	302.0	131.602649	17.563394	94.0	120.00	130.0	140.00	200.0
chol	302.0	246.500000	51.753489	126.0	211.00	240.5	274.75	564.0
fbs	302.0	0.149007	0.356686	0.0	0.00	0.0	0.00	1.0
restecg	302.0	0.526490	0.526027	0.0	0.00	1.0	1.00	2.0
thalach	302.0	149.569536	22.903527	71.0	133.25	152.5	166.00	202.0
exang	302.0	0.327815	0.470196	0.0	0.00	0.0	1.00	1.0
oldpeak	302.0	1.043046	1.161452	0.0	0.00	0.8	1.60	6.2
slope	302.0	1.397351	0.616274	0.0	1.00	1.0	2.00	2.0
ca	302.0	0.718543	1.006748	0.0	0.00	0.0	1.00	4.0
thal	302.0	2.324503	0.588366	1.0	2.00	2.0	3.00	3.0
target	302.0	0.543046	0.498970	0.0	0.00	1.0	1.00	1.0

5: Explore Categorical Variables

```

In [9]: # --- Explore Categorical Variables ---
# CORRECTED the variable names to match the actual columns in your data
categorical_vars = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal', 'target']
fig, axes = plt.subplots(3, 3, figsize=(18, 15))
fig.suptitle('Distribution of Categorical Variables', fontsize=20)

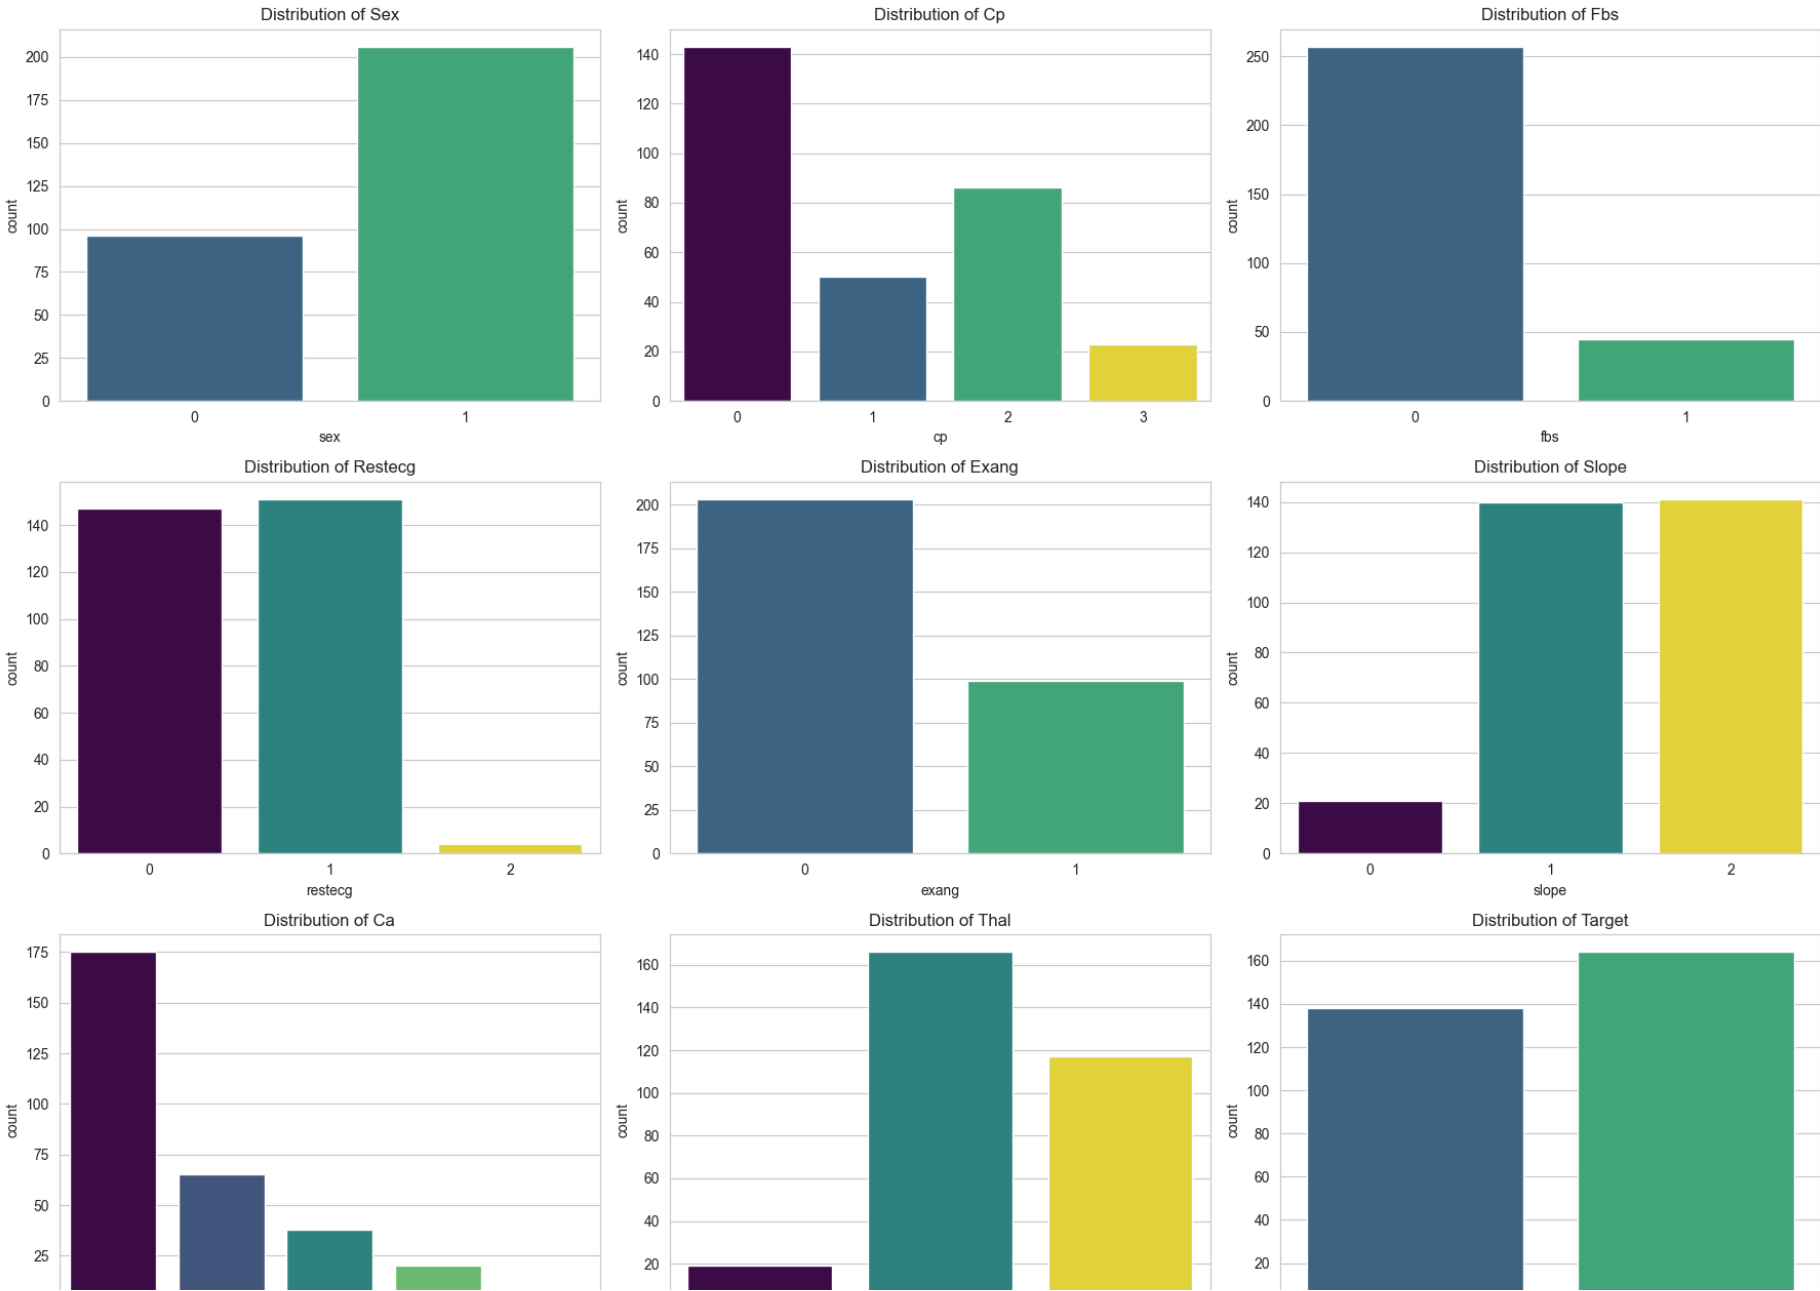
# Loop through the variables and create a count plot for each
for i, var in enumerate(categorical_vars):
    row, col = i // 3, i % 3

```

```
# FIX for FutureWarning: Assign the variable to 'hue' to use a palette and disable the legend
sns.countplot(x=var, data=df, ax=axes[row, col], hue=var, palette='viridis', legend=False)
axes[row, col].set_title(f'Distribution of {var.title()}')

plt.tight_layout(rect=[0, 0.03, 1, 0.95])
plt.savefig('../reports/figures/categorical_distributions.png', bbox_inches='tight')
plt.show()
```

Distribution of Categorical Variables






6: Occurrence of Heart Disease across Age

```
In [10]: # --- Resting Blood Pressure vs. Heart Disease ---
plt.figure(figsize=(10, 7))
# CORRECTED 'output' to 'target' and 'trtbps' to 'trestbps'
sns.boxplot(x='target', y='trestbps', data=df, palette='coolwarm')
plt.title('Resting Blood Pressure vs. Heart Disease Status', fontsize=16)
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Resting Blood Pressure (mm Hg)')
plt.xticks([0, 1], ['No Disease', 'Has Disease']) # More descriptive labels
plt.savefig('../reports/figures/bp_vs_heart_disease.png', bbox_inches='tight')
plt.show()
```

C:\Users\ThapeloMasebe\AppData\Local\Temp\ipykernel_52240\2457019085.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='target', y='trestbps', data=df, palette='coolwarm')
```



8: Patient Composition by Gender


```
In [11]: # --- Patient Composition by Gender ---
plt.figure(figsize=(8, 8))
df['sex'].value_counts().plot.pie(
    autopct='%1.1f%%', startangle=90, colors=['#6495ED', '#FFB6C1'],
    labels=['Male (1)', 'Female (0)'], textprops={'fontsize': 14},
    wedgeprops={'edgecolor': 'black', 'linewidth': 1}
)
plt.title('Overall Patient Composition by Gender', fontsize=16)
plt.ylabel('') # Hide the 'sex' label on the y-axis
plt.savefig('../reports/figures/gender_composition.png', bbox_inches='tight')
plt.show()
```

Overall Patient Composition by Gender

